


All seven test recordings: observed running and fixed-recipe predictions

Running above physical zero, measured in each mouse's training-target SD



Lines show means of three independently trained predictions; shading shows their range, not a confidence interval. No smoothing or window selection. Figures illustrate saved predictions; reported primary scores average individual-seed errors, not ensemble errors. Vertical scales differ by mouse. Native frames are not converted to seconds. These traces do not demonstrate forecasting, causality, or generated neural activity.